

Homework 1

*Student Name: Vũ Quang Sơn**UID: 23025086*

- Feel free to talk to other students in the class when doing the homework. You should, however, write down your solution yourself. You also must indicate on each homework with whom you collaborated and cite any other sources you use including Internet sites.
- You will write your solution in LaTeX and submit the pdf file in zip files, including relevant materials, through courses.uet.vnu.edu.vn
- Don't be late.

1 Homework 1 - 10pts

1.1 Let $f(x) = x^2 - x + 6$ and $\lambda = 0.05$.

First, calculate $\partial_x f$. Then starting from $x^{(0)} = 1 (k = 0)$, apply the gradient descent with $k = 1, 2, 3$.

We have first derivative of $f(x)$ is $\partial_x f = 2x - 1$
 Initializing $k = 0, x^{(0)} = 1 \Rightarrow f(x^{(0)}) = 6$

- After updating $k = 0$:

$$\partial_x f(x^{(0)}) = 2 * 1 - 1 = 1$$

$$x^{(1)} = 1 - 0.05 * 1 = 0.95$$

$$f(x^{(1)}) = 5.9525$$

- After updating $k = 1$:

$$\partial_x f(x^{(1)}) = 2 * 0.95 - 1 = 0.9$$

$$x^{(2)} = 0.95 - 0.05 * 0.9 = 0.905$$

$$f(x^{(2)}) = 5.914025$$

- After updating $k = 2$:

$$\partial_x f(x^{(2)}) = 2 * 0.905 - 1 = 0.81$$

$$x^{(3)} = 0.905 - 0.05 * 0.81 = 0.8645$$

$$f(x^{(3)}) = 5.88286025$$

- After updating $k = 3$:

$$\partial_x f(x^{(3)}) = 2 * 0.8645 - 1 = 0.729$$

$$x^{(4)} = 0.8645 - 0.05 * 0.729 = 0.82805$$

$$f(x^{(4)}) = 5.8576168025$$

2. Let $f(x) = (x_1 - x_2^2 - 3)^2 + (x_1 - x_2 - 1)^2$ and $\lambda = 0.2$. First, calculate $\partial_{x_i} f$. Then starting from $x^{(0)} = (-0.5, -0.5) (k = 0)$, apply the gradient descent with $k = 1, 2$.




Figure 1: $f(x)$ after every epoch

How many local maxima of f ? Could you find them?

The first derivative of $f(x)$:

- $\partial_{x_1} f = 2(x_1 - x_2^2 - 3) + 2(x_1 - x_2 - 1)$
- $\partial_{x_2} f = -4x_2(x_1 - x_2^2 - 3) - 2(x_1 - x_2 - 1)$

Initialing $x^{(0)} = (-0.5, -0.5)$, $k = 0$, $learning_rate = 0.2$, $f(x^{(0)})$

- After $k = 0$:

$$\begin{aligned}\partial_{x_1} f(x_0) &= -9.5 \\ \partial_{x_2} f(x_0) &= -5.5 \\ x^{(1)} &= [1.4, 0.6] \\ f(x^{(1)}) &= 15.0625\end{aligned}$$
- After $k = 1$:

$$\begin{aligned}\partial_{x_1} f(x_1) &= -4.32 \\ \partial_{x_2} f(x_1) &= 5.104 \\ x^{(2)} &= [2.264, -0.4208] \\ f(x^{(2)}) &= 3.672\end{aligned}$$
- After $k = 2$:

$$\begin{aligned}\partial_{x_1} f(x_2) &= 1.543 \\ \partial_{x_2} f(x_2) &= -4.906 \\ x^{(3)} &= [1.955, 0.560] \\ f(x^{(3)}) &= 2.002\end{aligned}$$

- After $k = 3$:

$$\partial_{x_1} f(x_3) = -1.928$$

$$\partial_{x_2} f(x_3) = 2.257$$

$$x^{(4)} = [2.341, 0.109]$$

$$f(x^{(4)}) = 1.968$$

To find the function's local maximum, we need to solve:

$$\frac{\partial f}{\partial x_1} = 0$$

$$\frac{\partial f}{\partial x_2} = 0$$

The result is: $x = [\frac{19}{8}, \frac{1}{2}]$ We have the Hessian matrix of $f(x)$:

$$A = \partial_{x_1 x_1} f(x_s)$$

$$B = \partial_{x_1 x_2} f(x_s)$$

$$C = \partial_{x_2 x_2} f(x_s)$$

Because $A * C - B^2 = 44 > 0$ and $A > 0 \Rightarrow f(x)$ has one local minimum at $x = [\frac{19}{8}, \frac{1}{2}]$

3. Apply the gradient descent with $\lambda = 0.01$ and $\lambda = 0.05$ respectively, and draw the contour plot of $f(x)$ after 10 steps

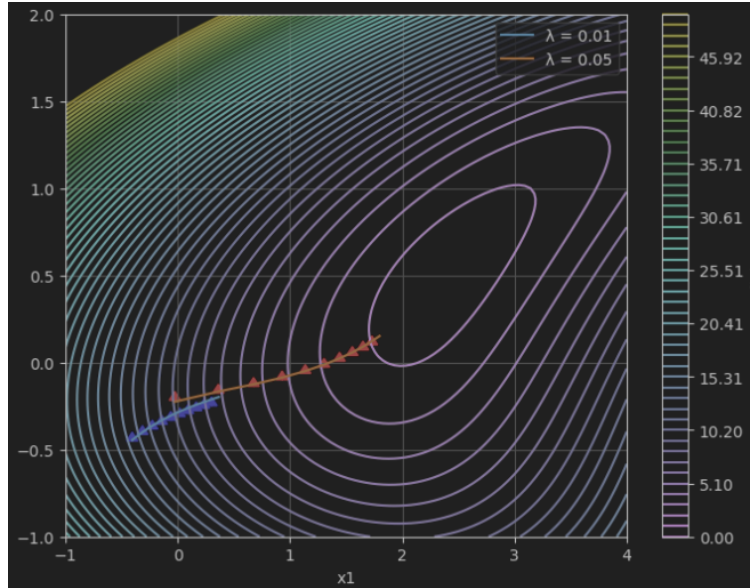


Figure 2: Gradient Descent with $\lambda = 0.01$ and $\lambda = 0.05$

Figure 2 shows that $\lambda = 0.05$ helps the function converge faster than $\lambda = 0.01$. Therefore, $\lambda = 0.05$ is the better for optimization.

2 Homework 2 - 10pts

Source code for generating data and model

```
def add_noise_data(input_data, input_labels, n_points, mean, scale):
    """
    Create a noise version of the input data

    Params:
        input_data: base input data
        input_labels: base input labels
        n_points: the number of needed points
        mean, scale: the Gaussian data
    """
    raw_X = []
    raw_labels = []

    noise = np.random.normal(loc=mean, scale=scale, size=(n_points, 2))
    for i in range(n_points):
        k = np.random.randint(len(input_data))

        raw_X.append([input_data[k][0] + noise[i][0],
                      input_data[k][1] + noise[i][1]])

        raw_labels.append(input_labels[k])

    return np.array(raw_X), np.array(raw_labels)

class LogisticRegression:
    def __init__(self, learning_rate=0.01, num_iterations=10000):
        self.learning_rate = learning_rate
        self.num_iterations = num_iterations
        self.weights = None
        self.bias = None

    def sigmoid(self, z):
        return 1 / (1 + np.exp(-z))

    def binary_cross_entropy_loss(self, y, y_pred):
        return -np.mean(y * np.log(y_pred) + (1 - y) * np.log(1 - y_pred))

    def error_calculation(self, y_pred, y_true):
        n = y_pred.shape[0]
        y_pred[np.where(y_pred >= 0.5)] = 1
        y_pred[np.where(y_pred < 0.5)] = 0
```

```

        return 1 / n * np.sum(np.square(y_pred - y_true))

def fit(self, X, y):
    n_samples, n_features = X.shape
    self.weights = np.zeros(n_features)
    self.bias = 0

    for i in range(self.num_iterations):
        linear_model = np.dot(X, self.weights) + self.bias
        y_predicted = self.sigmoid(linear_model)
        current_loss = self.binary_cross_entropy_loss(y, y_predicted)
        error = self.error_calculation(y_predicted, y)
        print(f"Epoch {i + 1}: Loss = {current_loss}, weights = {self.weights}, bias = {self.bias}")
        if error == 0:
            print(f"Done")
            break

        dw = (1 / n_samples) * np.dot(X.T, (y_predicted - y))
        db = (1 / n_samples) * np.sum(y_predicted - y)

        self.weights -= self.learning_rate * dw
        self.bias -= self.learning_rate * db

def predict(self, X):
    linear_model = np.dot(X, self.weights) + self.bias
    y_predicted = self.sigmoid(linear_model)
    y_predicted_cls = [1 if i > 0.5 else 0 for i in y_predicted]
    return y_predicted_cls

```

2.1 We have logistic function h which satisfies the condition $er\hat{r}_D(h) = 0$:

$$\begin{aligned}
 P\{|er\hat{r}_D(h) - err_{\mathbb{P}}(h)| \leq \epsilon\} &\geq 1 - 2e^{-2n\epsilon^2} \\
 \Leftrightarrow P\{|-err_{\mathbb{P}}(h)| \leq \epsilon\} &\geq 1 - 2e^{-2n\epsilon^2} \\
 \Leftrightarrow P\{err_{\mathbb{P}}(h) \leq \epsilon\} &\geq 1 - 2e^{-2n\epsilon^2}
 \end{aligned}$$

With $\epsilon = 0.1, n = 10$:

$$\Rightarrow P\{err_{\mathbb{P}}(h) \leq 0.1\} \geq 1 - 2e^{-2*10*0.1^2} = -0.637$$

The above shows that $er\hat{r}_D(h) = 0$ with $P\{err_{\mathbb{P}}(h) > 0.1 = 0\}$ or we can say this does not happen.

2.2 `X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])`
`y = np.array([0, 0, 0, 1])`
`seeds = [0, 10, 20, 30, 40]`

```

for seed in seeds:
    print(f"Seed {seed}")
    np.random.seed(seed)
    X_noise, y_noise = add_noise_data(X, y, 10, 0., 0.2)
    model = LogisticRegression()
    model.fit(X_noise, y_noise)

```

Seed	Number of Epoches
0	8
10	16
20	9
30	14
40	6

Table 1: The number of epochs to converge respectively to seed

2.3 When the sample size increases, the time for the algorithm to classify increases and can not be separable.

```

X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = np.array([0, 0, 0, 1])
seeds = [0, 10, 20, 30, 40, 50]
N = [20, 50, 100, 200, 500]
for n in N:
    print("N = {0}".format(n))
    for seed in seeds:
        print(f"Seed {seed}")
        np.random.seed(seed)
        X_noise, y_noise = add_noise_data(X, y, n, 0., 0.2)
        model = LogisticRegression()
        model.fit(X_noise, y_noise)

```

Samples	Seed	Number of Epochs
20	0	10
	10	21
	20	5
	30	9
	40	8
50	0	26
	10	11
	20	Failed
	30	10
	40	11
100	0	Failed
	10	Failed
	20	Failed
	30	Failed
	40	Failed
200	0	Failed
	10	Failed
	20	Failed
	30	Failed
	40	Failed
500	0	Failed
	10	Failed
	20	Failed
	30	Failed
	40	Failed

Table 2: The number of epochs to converge respectively to samples, seed

```

2.4 X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = np.array([0, 0, 0, 1])
seeds = [0, 42, 100, 200]
N = [10, 20]
for n in N:
    print("N = {0}".format(n))
    for seed in seeds:
        print(f"Seed {seed}")
        np.random.seed(seed)
        X_noise, y_noise = add_noise_data(X, y, n, 0., 0.2)
        X_test, y_test = add_noise_data(X, y, n, 0., 0.3)
        model = LogisticRegression()
        model.fit(X_noise, y_noise)
        y_pred = model.sigmoid(np.dot(X_test, model.weights) + model.bias)
        print(f"Testing ERR: {model.error_calculation(y_test, y_pred)}")

```

Samples	Seed	Number of Epoches	Training Error	Testing Error
10	0	8	0	0.2492
	42	19	0	0.2487
	100	6	0	0.2495
	200	10	0	0.2493
20	0	10	0	0.2493
	42	9	0	0.2494
	100	12	0	0.2495
	200	11	0	0.2492

Table 3: Result of 2.3